

# Ejercicios - Algoritmos Paralelos

Lorena Xiomara Castillo Galdos

Junio 2017

- Ejercicios cuda

- Suma de matrices

```
#include <iostream>
#include <cuda.h>
using namespace std;

__global__ void MatrizAddKernel_A(float *A, float *B, float *C,
    int n)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if(i < n*n)
    {
        C[i] = A[i] + B[i];
    }
}

__global__ MatrizAddKernel_B(float *A, float* B, float* C,
    int n)
{
    int i=threadIdx.x+blockDim.x*blockIdx.x;
    if(i < n*n)
    {
        C[i]=A[i]+B[i];
    }
}

__global__ MatrizAddKernel_C(float *A, float* B, float* C,
    int n)
{
    int i= n*blockIdx.x;
    if(i < n*n)
    {
        for(int j = 0 ; j < n ; j++)
        {
            C[i + j]=A[i + j]+B[i + j];
        }
    }
}

__global__ MatrizAddKernel_D(float *A, float* B, float* C,
    int n)
```

```

{
    int i= blockIdx.x;
    if(i < n*n)
    {
        for(int j = 0 ; j < n ; j++)
        {
            C[i+n*j]=A[i+n*j]+B[i+n*j];
        }
    }
}

void MatrizAdd(float *A, float *B, float *C, int n)
{
    float *d_A, *d_B, *d_C;
    size_t size = n*n * sizeof(float);

    cudaMalloc((void **) &d_A, size);
    cudaMalloc((void **) &d_B, size);
    cudaMalloc((void **) &d_C, size);

    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    //MatrizAddKernel_A<<< 5 , 2 >>>(d_A, d_B, d_C, n);
    MatrizAddKernel_B<<< 1 , n*n >>>(d_A, d_B, d_C, n);
    //MatrizAddKernel_C<<< n , n >>>(d_A, d_B, d_C, n);
    //MatrizAddKernel_D<<< n , n >>>(d_A, d_B, d_C, n);

    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}

int main(void)
{
    float *A, *B , *C;
    int n = 10 ;
    A = new float [n*n];
    B = new float [n*n];
    C = new float [n*n];
    for (int i = 0; i < n*n; ++i)
    {
        A[i] = i;
        B[i] = i;
    }

    MatrizAdd(A, B, C ,n);

    for(int i = 0; i < n*n ; ++i)
    {
        cout << C[i] << " ";
    }
    return 0;
}

```

- Multiplicación de matriz vector

```
#include <iostream>
#include <cuda.h>
using namespace std;

__global__ void MultMatrizVectKernel(float *A, float *B, float *C,
                                     int n)
{
    int i = n * blockIdx.x;
    if(i < n*n)
    {
        for(int j = 0; j < n ; ++j)
        {
            C[blockIdx.x] += A[i + j] * B[j];
        }
    }
}

void MultMatrizVector(float *A, float *B, float *C, int n)
{
    float *d_A, *d_B, *d_C;
    size_t size_A = n*n * sizeof(float);
    size_t size_B = n * sizeof(float);
    size_t size_C = n * sizeof(float);

    cudaMalloc((void **) &d_A, size_A);
    cudaMalloc((void **) &d_B, size_B);
    cudaMalloc((void **) &d_C, size_C);

    cudaMemcpy(d_A, A, size_A, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, B, size_B, cudaMemcpyHostToDevice);

    MultMatrizVectKernel<<<< n, 1 >>>>(d_A, d_B, d_C, n);

    cudaMemcpy(C, d_C, size_C, cudaMemcpyDeviceToHost);

    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}

int main(void)
{
    float *A, *B, *C;
    int n = 10 ;
    A = new float[n*n];
    B = new float[n];
    C = new float[n];
    for (int i = 0; i < n*n; ++i)
    {
        A[i] = i;
    }

    for (int i = 0; i < n; ++i)
    {
        B[i] = i;
    }

    MultMatrizVector(A, B, C, n);
}
```

```

    for(int i = 0; i < n ; ++i)
    {
        cout << C[i] << " ";
    }
    return 0;
}

```

- Imágenes blurr y gray

```

#include <cv.h>
#include <highgui.h>
#include <iostream>
#include <math.h>

using namespace std;

#define CHANNELS 3

__global__ void colorConvertKernel(unsigned char * greyImage,
    unsigned char * rgbImage, int width, int height)
{
    int Col = threadIdx.x + blockIdx.x * blockDim.x;
    int Row = threadIdx.y + blockIdx.y * blockDim.y;

    if (Col < width && Row < height)
    {
        int greyOffset = Row*width + Col;
        int rgbOffset = greyOffset*CHANNELS;
        unsigned char r = rgbImage[rgbOffset];
        unsigned char g = rgbImage[rgbOffset + 2];
        unsigned char b = rgbImage[rgbOffset + 3];
        greyImage[grayOffset] = 0.21f*r + 0.71f*g + 0.07f*b;
    }
}

//1px
__global__ void blurrKernel(unsigned char * in,
    unsigned char * out, int w, int h)
{
    int Col = blockIdx.x * blockDim.x + threadIdx.x;
    int Row = blockIdx.y * blockDim.y + threadIdx.y;

    if (Col < w && Row < h)
    {
        int pixVal = 0;
        int pixels = 0;

        //promedio de blur_size x blur_size box
        for(int blurRow = -BLUR_SIZE; blurRow < BLUR_SIZE+1;
            ++blurRow)
        {
            for(int blurCol = -BLUR_SIZE; blurCol < BLUR_SIZE+1;
                ++blurCol)
            {
                int curRow = Row + blurRow;

```

```

        int curCol = Col + blurCol;
//verifica pixel
        if (curRow > -1 && curRow < h && curCol > -1
            && curCol < w)
        {
            pixVal += in[curRow * w + curCol];
            pixels++;
        }
    }
    out[Row * w + Col] = (unsigned char)(pixVal / pixels);
}

int main(int argc, char** argv)
{
    IplImage* input_image = NULL; //datos imagen
    input_image = cvLoadImage(argv[1], CV_LOAD_IMAGE_UNCHANGED);

    int width = input_image->width;
    int height = input_image->height;

    float* imagen_cpu = new float[width * height * 4];
    float* imagen_gpu = new float[width * height * 4];

    cudaMalloc((void **>(&imagen_gpu), (width * height * 4)
        * sizeof(float));
    cudaMemcpy(imagen_gpu, imagen_cpu, (width * height * 4)
        * sizeof(float), cudaMemcpyHostToDevice);

//16Thrd
    dim3 dimGrid(ceil(width/16.0), ceil(height/16.0), 1);
    dim3 dimBlock(16, 16, 1);

    colorConvertKernel<<<dimGrid,dimBlock>>>(imagen_gpu, input_image,
        width, height);

    blurrkernel<<<dimGrid,dimBlock>>>(imagen_gpu, input_image,
        width, height)

    cudaMemcpy(imagen_cpu, imagen_gpu, (width * height * 4)
        * sizeof(float), cudaMemcpyDeviceToHost);
    cvReleaseImage(&input_image);
    cvReleaseImage(&out_image);
    return 0;
}

```

- **Openmp**

- Odd even sort

```

#include <iostream>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>

```

```

#include <chrono>
using namespace std;

void odd_even(int *V, int n, int thread_count)
{
    int fase,i,tmp;
    for (fase = 0; fase < n; fase++)
    {
        if (fase % 2 == 0)
            #pragma omp parallel for num_threads(thread_count) \
            default(none) shared(a, n) private(i, tmp)
            for (i = 1; i < n; i += 2)
            {
                if (V[i - 1] > V[i])
                {
                    tmp = V[i - 1];
                    V[i - 1] = V[i];
                    V[i] = tmp;
                }
            }
        else
            #pragma omp parallel for num_threads(thread_count) \
            default(none) shared(a, n) private(i, tmp)
            for (i = 1; i < n - 1; i += 2)
            {
                if (V[i] > V[i+1])
                {
                    tmp = V[i+1];
                    V[i+1] = V[i];
                    V[i] = tmp;
                }
            }
    }
}

void odd_even_(int *V, int n, int thread_count)
{
    int fase,i,tmp;
    #pragma omp parallel num_threads(thread_count) \
    default(none) shared(a, n) private(i, tmp, fase)
    for (fase = 0; fase < n; fase++)
    {
        if (fase % 2 == 0)
            #pragma omp for
            for (i = 1; i < n; i += 2)
            {
                if (V[i - 1] > V[i])
                {
                    tmp = V[i - 1];
                    V[i - 1] = V[i];
                    V[i] = tmp;
                }
            }
        else
            #pragma omp for
            for (i = 1; i < n - 1; i += 2)
            {

```

```

        if (a[i] > a[i+1])
        {
            tmp = a[i+1];
            a[i+1] = a[i];
            a[i] = tmp;
        }
    }
}

void print(int *V,int n)
{
    for (int i = 0; i < n; ++i)
    {
        cout<<V[i]<<" ";
    }
    cout<<endl;
}

int main(int argc , char const *argv [])
{
    int thread_count=4;
    int n=20000;
    int *V=new int [n];
    srand (time(NULL));
    for (int i = 0; i < n; ++i)
    {
        int iSecret = rand() % n;
        V[i]=iSecret;
    }
    auto start = chrono::high_resolution_clock::now();
    odd_even(V,n,thread_count);

    auto finish = chrono::high_resolution_clock::now();
    chrono::duration<double> elapsed = finish - start;
    cout << "Time_taken_First_Odd-Even:" << elapsed.count() <<endl;

    auto start2 = chrono::high_resolution_clock::now();
    odd_even_(V,n,thread_count);

    auto finish2 = chrono::high_resolution_clock::now();
    chrono::duration<double> elapsed2 = finish2 - start2;
    cout << "Time_taken_Second_Odd-Even:" << elapsed2.count() <<endl;
    return 0;
}

```